

# Docker for Beginners

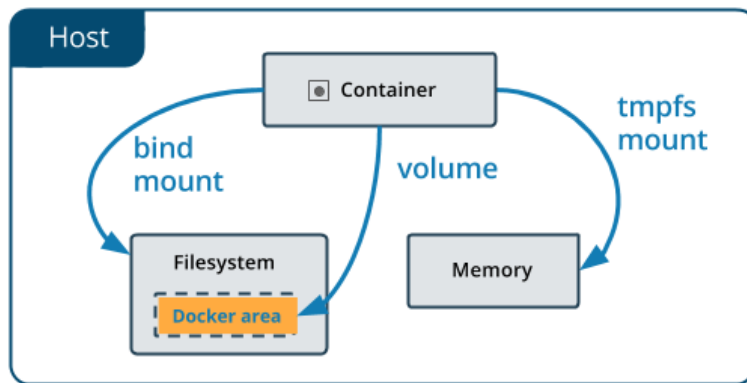
## Day 4 – Volumes and Networking

- Lucas Albuquerque -  
lucas.albuquerque@nutanix.com

# > Docker Volumes

- Containers are **ephemeral** and mere processing units **by definition**
- Docker container is **destroyed**, any stored data will be **lost**.
- Docker **volumes** are very useful when we need to **persist data** in Docker containers or share data between containers.
- There are **three ways** to keep data safe in case of a container removal:

- ✓ Bind Mounts
- ✓ Named Volumes
- ✓ tmpfs



# > Docker Volumes

## Bind Mounts

- A bind mount is just **mapping a directory** on the host machine to a directory in the container.
- However, when the container is **removed**, it does **not affect the directory**.
- Content can be **modified** by other applications

```
$ docker run -d --name nginx-app \  
-v /path/to/app/directory:/usr/share/nginx/html \  
hutger/nginx-app:0.1
```

# > Docker Volumes

## Named Volumes

- Volumes which you create manually with docker ***volume create VOLUME\_NAME***.
- Created in `/var/lib/docker/volumes` and can be referenced to by only their name.
- It is not an Union FS storage entity;
- Should not be modified by other applications;

```
$ docker volume create nginx_vol
$ docker volume ls

$ docker run -d --name nginx-app \
-v nginx_vol:/usr/share/nginx/html \
hutger/nginx-app:0.1
```

# > Docker Volumes

## Named Volumes

- You can create or connect to **remote volumes** using **backend** drivers:

```
$ docker volume create --driver local \  
  --opt type=nfs \  
  --opt o=addr=192.168.1.1,rw \  
  --opt device=:/path/to/dir \  
  nginx-vol
```

- Backend driver options by adding plugins: **Nutanix DVP, Azure, Flocker, DRDB, GCE, GlusterFS, 3Par, NetApp, vSphere, etc.**

[https://docs.docker.com/engine/extend/legacy\\_plugins/](https://docs.docker.com/engine/extend/legacy_plugins/)

# > Docker Volumes

## tmpfs

- If you do **not want the data to persist** either on the host or container
- For **security** reasons or to protect the **performance** of the container
- tmpfs mount is **temporary**, and only persisted in the **host memory**
- Unlike volumes and bind mounts, you **can't share** tmpfs mounts
- This functionality is **only available** if you're running Docker on **Linux**

```
$ docker run -d --name nginx-app \  
--tmpfs /usr/share/nginx/html/uploads \  
hutger/nginx-app:0.1
```

# > Docker Networking

## Linux Bridge

- A **Linux bridge** is a L2 device that is the **virtual** implementation of a network **switch** inside the Linux kernel.

## Network Namespaces

- **Isolated** network stack in the kernel with its **own interfaces, routes, and firewall rules**.
- It is a security aspect of containers and Linux, used to **isolate containers**.

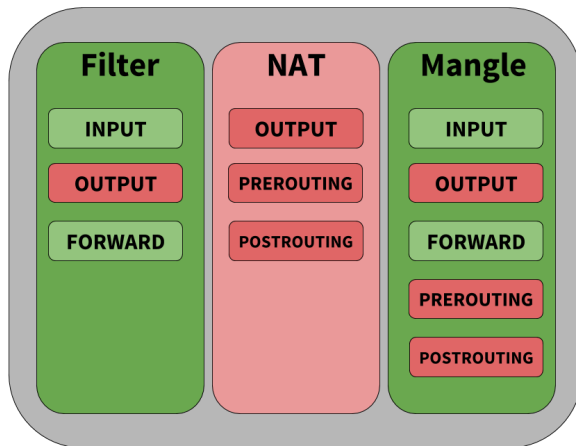
## Virtual Ethernet Devices

- A **veth** is a full duplex link that has a **single interface** in each namespace.
- Docker network drivers utilize veths to provide **explicit connections** between namespaces when Docker networks are created.

# > Docker Networking

## IPtables

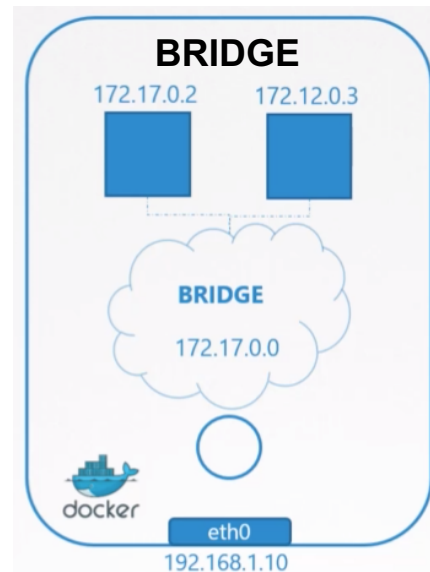


- It's a feature rich L3/L4 firewall that provides rule chains for packet marking, masquerading, and dropping.
- Docker network drivers utilize iptables extensively to segment network traffic, provide host port mapping, and to mark traffic for load balancing decisions.





# > Docker Networking

```
vagrant@ubuntu-xenial:~/docker-training$ docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
4bb793eb33ff        bridge             bridge              local
0de98930f170        host               host                local
9bcf1c3f912b        none               null                local
```



# > Docker Networking


```
vagrant@ubuntu-xenial:~/docker-training$ ip addr | grep docker
3: docker0: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc noqueue state UP group default
   inet 172.17.0.1/16 brd 172.17.255.255 scope global docker0
```

- Everytime a new container is create, a new network space is also created:

```
root@ubuntu-xenial:~# docker inspect d6c2cf16977d | grep -i Networksettings -A10
  "NetworkSettings": {
    "Bridge": "",
    "SandboxID": "42b0f49974768e396e4462eaa1c5d5a825f59538595f481302a06025f827dfd4",
    "HairpinMode": false,
    "LinkLocalIPv6Address": "",
    "LinkLocalIPv6PrefixLen": 0,
    "Ports": {
      "80/tcp": null
    },
    "SandboxKey": "/var/run/docker/netns/42b0f4997476",
    "SecondaryIPAddresses": null,
```

# > Docker Networking

- Each container is connected to **bridge docker0** using a **veth**;
- Each veth connects to **eth0** on container and to **docker0** on host (vethxxxx);
- **Netfilter** running on host is responsible for all **NATs and REDIRECTs** from/to containers;



# > Docker Networking

## Linking Container

- Docker **doesn't** provide an native **naming service** solution
- In order to make reference to a **different container** by name, you can use the **resource --link** while running a container:

```
$ docker run --name mariadb-ping-test --rm \
--link my-mariadb \
-it busybox sh -c "ping -c5 my-mariadb"
PING my-mariadb (172.17.0.4): 56 data bytes
64 bytes from 172.17.0.4: seq=0 ttl=64 time=0.188 ms
64 bytes from 172.17.0.4: seq=1 ttl=64 time=0.170 ms
64 bytes from 172.17.0.4: seq=2 ttl=64 time=0.105 ms
64 bytes from 172.17.0.4: seq=3 ttl=64 time=0.137 ms
64 bytes from 172.17.0.4: seq=4 ttl=64 time=0.104 ms

--- my-mariadb ping statistics ---
5 packets transmitted, 5 packets received, 0% packet loss
round-trip min/avg/max = 0.104/0.140/0.188 ms
```

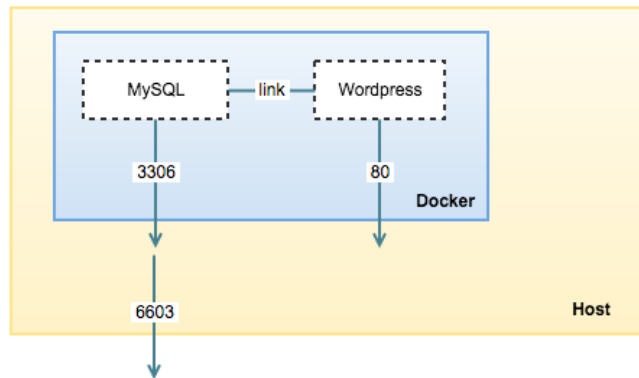
- The magic is made by basically **adding** an new entry on **/etc/hosts** of the running container with the link information.

172.17.0.4 my-mariadb
-----------------------

# > Docker Networking

## Publishing Ports

- By default, the container is **assigned an IP** for every Docker network it connects to.
- The IP is assigned **from the pool** assigned to the network. Docker daemon **acts as a DHCP server**;
- **Containers** inside of the **same host can communicate** to each other (as long as the IPs are known)
- **Publish** service is a way to **allow external users** to connect to container services;



# > Docker Networking

## Publishing Ports

```
# It will make a bind from port 80/TCP on host to port 80/TCP on the container
1 $ docker run -d -p 80:80 nginx

# It will make a bind from port 8080/TCP on host to port 80/TCP on the container
2 $ docker run -d -p 8080:80 nginx

# Publish all exposed ports to random ports. It relies on EXPOSE settings in Dockerfile.
3 $ docker run -d -P nginx
```

```
root@ubuntu-xenial:~# docker ps
```

	CONTAINER ID	IMAGE	COMMAND	CREATED	STATUS	PORTS
3	bb5371424fa9	nginx	"nginx -g 'daemon of..."	4 seconds ago	Up 3 seconds	0.0.0.0:32770->80/tcp
2	cba2a858af5b	nginx	"nginx -g 'daemon of..."	8 seconds ago	Up 7 seconds	0.0.0.0:8080->80/tcp
1	ad4fef079676	nginx	"nginx -g 'daemon of..."	18 seconds ago	Up 18 seconds	0.0.0.0:80->80/tcp

```
root@ubuntu-xenial:~# iptables -t nat -L DOCKER
```

	target	prot	opt	source	destination
	RETURN	all	--	anywhere	anywhere
1	DNAT	tcp	--	anywhere	anywhere tcp dpt:5000 to:172.17.0.4:5000
2	DNAT	tcp	--	anywhere	anywhere tcp dpt:32769 to:172.17.0.3:5000
3	DNAT	tcp	--	anywhere	anywhere tcp dpt:http to:172.17.0.6:80
	DNAT	tcp	--	anywhere	anywhere tcp dpt:http-alt to:172.17.0.7:80
	DNAT	tcp	--	anywhere	anywhere tcp dpt:32770 to:172.17.0.8:80

# > Docker Volumes and Networking





# See you tomorrow...



- Lucas Albuquerque -  
lucas.albuquerque@nutanix.com





Thank You